

COURSE ASSIGNMENT**CSA04 - OPERATING SYSTEMS****Process Management, Memory Management & Disk Scheduling****A. Assignment Information**

Department	Computer Science and Engineering
Programme	B.E. CSE
Course Code & Course Name	CSA04 — Operating Systems
Academic Year / Batch	2026 – 2027
Faculty Name	Dr. Sundar Rajan S
Assignment Title	Design and Implementation of CPU Scheduling, Memory Management and Disk Scheduling Strategies for a Real-Time Multi-Service System
Assignment Type	Implementation-based, Team Assignment (Team size: 3 students)
Date of Issue	
Date of Submission	02/09/2026
Maximum Marks	100
Course Outcome(s) – CO	CO2, CO3, CO4
Bloom's Taxonomy Level	L3 – Apply, L4 – Analyze, L5 – Evaluate
SDG Mapping (SDG 1–17)	SDG 9 – Industry, Innovation and Infrastructure
Industry / Societal Relevance	Efficient CPU scheduling, memory management, and disk I/O scheduling are foundational to the performance of servers, cloud platforms, and real-time systems used across every industry sector.

1. Problem Statement and Problem Formulation

An organization runs a Linux-based Library Management System (LMS) server that must simultaneously handle multiple concurrent user activities on a single CPU core — catalog search, book issue/return transactions, member-record updates, fine calculation, and report generation. The server must (a) schedule these request-handling processes fairly and efficiently on the CPU, (b) manage a limited pool of physical RAM across many simultaneous process address spaces using virtual memory, and (c) schedule disk I/O requests (catalog database reads/writes, receipt printing spools, backup writes) efficiently to minimize seek time. This assignment restates that scenario for assigned parameter set and designs, implements and validates an integrated OS resource-management scheme covering CPU scheduling, memory management, and disk scheduling for this domain.

In the LMS domain, the specific challenges are: (i) request bursts during peak hours (start of semester, exam-time renewals) create CPU contention among catalog-search, issue/return and report-generation processes of different lengths; (ii) each active session (librarian terminal or self-service kiosk) occupies a logical address space that must be mapped into a physically limited 4 GB RAM pool via paging, with page faults costing time; and (iii) every transaction ultimately reads/writes the catalog database and log files on disk, so head movement must be minimised to keep response times low.

1.1 Objective and Expected Outcomes

- Model the LMS server's incoming requests as a process set and evaluate FCFS, SJF, Priority and Round-Robin CPU scheduling to identify the algorithm that minimises average waiting/turnaround time while remaining fair.
- Design the paging scheme for a representative LMS session process and compare FIFO, LRU and Optimal page-replacement policies for the assigned frame allocation.
- Compute and compare total head movement for FCFS, SSTF, SCAN and C-SCAN disk-scheduling algorithms for the LMS database/catalog I/O queue.
- Produce a single, quantitatively justified, integrated OS resource-management design for the LMS server, validated by simulation.
- Expected outcome: a working Python simulation, comparison tables/graphs for all three subsystems, and a recommended configuration (scheduling algorithm + paging policy + disk algorithm) suited to the LMS workload.

1.2 Requirements, Constraints and Assumptions

Requirements (from assignment Section D)

- Functional: CPU scheduling must minimise average waiting/turnaround time while remaining fair to all processes.
- Performance: Average response time must stay low under peak concurrent load; disk seek overhead must be minimised.
- Technical: Physical RAM = 4 GB, disk cylinders = 0–130
- Reliability: No process may be starved indefinitely under the chosen CPU/disk scheduling policy.
- Cost & Time: Implementation restricted to open-source tools (Python); completed within the assignment timeline.
- Standards: Follow standard OS design conventions as covered in Units II–IV of the syllabus.

Constraints

- CPU process set: P1(0,6), P2(2,2), P3(3,5), P4(5,9), P5(7,3) — (Arrival Time, Burst Time) in ms; Round-Robin quantum = 4 ms.
- Memory: RAM = 4 GB, Page size = 4 KB, Process logical address space = 32 MB, Frames allocated to the representative process = 4.
- Disk: Cylinders 0–130, initial head position = 65, request queue = 25, 105, 40, 115, 55, 90, 10, 120.

Assumptions

- Process burst times are known/estimated in advance for non-preemptive SJF, as is conventional for this exercise; in a real LMS server these would be estimated using exponential averaging of past bursts.
- Priorities (1 = highest) are assigned by shortest-burst-first-served heuristic reflecting business urgency: P2 (fast catalog lookup) = 1, P1 (issue transaction) = 2, P3 (return + fine calc) = 3, P5 (member update) = 4, P4 (bulk report generation) = 5 — justified further in Section 6.
- All five processes are CPU-bound with no I/O-wait sub-phases for the CPU-scheduling simulation (I/O is modelled separately in the disk-scheduling part).
- The representative process for memory management is a librarian session (search → view → issue → update → print) occupying the full 32 MB logical space with 4 physical frames allocated (as given), reflecting realistic frame contention under multi-user load.
- Disk head movement is measured in cylinder units on a linear disk model, ignoring rotational latency, consistent with standard textbook treatment of disk-scheduling algorithms.

2. Application of Relevant Course Knowledge / Concepts

This section answers Part I–III of Section E.2 of the assignment using parameter set, with full manual working shown alongside the simulated results (Section 5 in this report reproduces the same figures from the Python implementation for cross-validation).

2.1 Part I — Process Management & CPU Scheduling (CO2, Unit II)

Process set: P1(0,6), P2(2,2), P3(3,5), P4(5,9), P5(7,3). All times in milliseconds.




Figure 1: Process state-transition diagram (new → ready → running → waiting/terminated) relevant to the LMS request-handling scenario.

(a) FCFS — First Come First Served

Processes are serviced strictly in arrival order: P1, P2, P3, P4, P5.

PID	AT	BT	CT	TAT	WT
P1	0	6	6	6	0
P2	2	2	8	6	4
P3	3	5	13	10	5
P4	5	9	22	17	8
P5	7	3	25	18	15

Average Waiting Time = 6.40 ms | Average Turnaround Time = 11.40 ms

(b) SJF — Shortest Job First (Non-preemptive)

At each dispatch point, the ready process with the smallest burst time is chosen: order P1, P2, P5, P3, P4.

PID	AT	BT	CT	TAT	WT
P1	0	6	6	6	0
P2	2	2	8	6	4
P5	7	3	11	4	1
P3	3	5	16	13	8
P4	5	9	25	20	11

Average Waiting Time = 4.80 ms | Average Turnaround Time = 9.80 ms

(c) Priority Scheduling (Non-preemptive, 1 = highest priority)

Assigned priorities: P1 = 2, P2 = 1, P3 = 3, P4 = 5, P5 = 4 (justification in Section 1.2 / Section 6). Order: P1, P2, P3, P5, P4.

PID	Priority	AT	BT	CT	TAT	WT
P1	2	0	6	6	6	0
P2	1	2	2	8	6	4
P3	3	3	5	13	10	5
P5	4	7	3	16	9	6
P4	5	5	9	25	20	11

Average Waiting Time = 5.20 ms | Average Turnaround Time = 10.20 ms

(d) Round Robin (Quantum = 4 ms)

Execution order (Gantt sequence): P1(0-4), P2(4-6), P3(6-10), P1(10-12), P4(12-16), P5(16-19), P3(19-20), P4(20-24), P4(24-25).

PID	AT	BT	CT	TAT	WT
P1	0	6	12	12	6
P2	2	2	6	4	2
P3	3	5	20	17	12
P4	5	9	25	20	11
P5	7	3	19	12	9

Average Waiting Time = 8.00 ms | Average Turnaround Time = 13.00 ms




Figure 2: Gantt charts for FCFS, SJF, Priority and Round Robin (simulated).

Comparison and Recommendation

Algorithm	Avg. Waiting Time (ms)	Avg. Turnaround Time (ms)
FCFS	6.40	11.40
SJF (Non-preemptive)	4.80	9.80
Priority (Non-preemptive)	5.20	10.20
Round Robin (q=4)	8.00	13.00

SJF gives the lowest average waiting and turnaround time for this process set, but it is not fair — a long process (P4, burst 9) can be delayed behind a stream of short ones, and true burst times are rarely known in advance on a live server. Round Robin has the highest average metrics here because of the short quantum relative to burst sizes, but it bounds the maximum wait any process experiences and guarantees fairness, which is required by Section D ("remaining fair to all processes"; "must not starve any process"). For the LMS domain — a mix of quick catalog look-ups and long report jobs, serving many simultaneous library users — Round Robin ($q = 4$) is recommended as the primary scheduler because it prevents any single long-running report job from starving interactive catalog/issue requests, at an acceptable cost in average turnaround time. Priority scheduling is retained as a secondary policy for distinguishing latency-sensitive lookups when the system is lightly loaded.

2.2 Part II — Memory Management (CO3, Unit III)

Physical RAM = 4 GB = 2^{32} bytes. Page size = 4 KB = 2^{12} bytes.

- Total physical frames in system RAM = $2^{32} / 2^{12} = 2^{20} = 1,048,576$ frames.

- Representative LMS session process logical address space = 32 MB = 2^{25} bytes.
- Pages required for this process = $2^{25} / 2^{12} = 2^{13} = 8,192$ pages.
- Page-table structure: 8,192 entries, one per logical page; with a 20-bit frame number and typical control bits (valid/invalid, dirty, referenced, protection) each entry is a few bytes, giving a single-level page table of a few tens of KB per process — acceptable, but on a heavily loaded LMS server with many concurrent sessions this motivates a multi-level or inverted page table to bound page-table memory overhead.
- Logical address breakdown: 25-bit logical address = 13-bit page number + 12-bit page offset.

Frames allocated to this representative process = 4 . A 12-reference page-access string was constructed to reflect a typical LMS user session — repeated navigation between the catalogue-search page, book-detail page, cart/issue page and confirmation page, with an excursion to the fine-payment page:

Reference string (logical page numbers): 0, 1, 2, 3, 0, 1, 4, 0, 1, 2, 3, 4

Page-Replacement Comparison (4 frames)

Algorithm	Page Faults	Fault Rate
FIFO	10	83.3%
LRU	8	66.7%
Optimal	6	50.0%




Figure 3: Page-fault count comparison — FIFO vs LRU vs Optimal (4 frames, 12 references).

Optimal is the theoretical lower bound and is not implementable in a live system since it needs future knowledge; LRU tracks recency and, on this reference string, comes close to Optimal (8 vs 6 faults) because the LMS access pattern shows temporal locality (the user revisits the same 3-4 pages repeatedly). FIFO performs worst because it evicts the oldest-loaded page even if it was just re-referenced, which happens often here (page 0 and 1 are re-used but FIFO still cycles them out).

Demand Paging and the Working-Set Model

Demand paging loads a page into a frame only on first reference (page fault), avoiding wasted RAM for pages of a session the user never visits (e.g., admin-only report pages within a librarian's session). Combined with the working-

set model — tracking the set of pages referenced in the last Δ references — the LMS server can detect when a session's working set (here, effectively 4 pages: search, detail, cart, confirmation) exceeds its allocated frames and is thrashing. Given the reference string above, a working-set window of $\Delta = 5$ stabilises at approximately 3–4 pages, matching the 4-frame allocation; this supports a recommended per-session frame allocation of 4–5 frames with LRU replacement, and a global working-set-based admission policy that suspends new sessions rather than over-committing frames when aggregate working-set size approaches the system's total frame budget under peak load.

2.3 Part III — Disk Scheduling (CO4, Unit IV)

File Allocation Method

Method	Suitability for LMS catalog/records
Contiguous	Fast sequential + direct access, ideal for read-mostly catalog index files, but suffers external fragmentation as records grow — poor fit for a frequently updated member/transaction database.
Linked	No external fragmentation and easy growth, but only efficient for sequential access; random lookups (member ID search) require traversing the chain — too slow for LMS's frequent random reads.
Indexed	Supports fast random and sequential access via an index block per file, and files can grow without fragmentation — best match for the LMS catalog/member database's mix of random lookups and updates.

Recommendation: Indexed allocation for the catalog/member/transaction database files (frequent random access + growth); contiguous allocation for read-only/append-only files such as daily backup or report archives where sequential access dominates.

Disk-Scheduling Comparison

Cylinders 0–130, head start = 65, request queue = 25, 105, 40, 115, 55, 90, 10, 120.

Algorithm	Service Order	Total Head Movement (cylinders)
FCFS	25, 105, 40, 115, 55, 90, 10, 120	545
SSTF	55, 40, 25, 10, 90, 105, 115, 120	165
SCAN (moving toward 130 first)	90, 105, 115, 120, 130, 55, 40, 25, 10	185
C-SCAN (moving toward 130 first)	90, 105, 115, 120, 130, 0, 10, 25, 40, 55	250




Figure 4: Total head movement — FCFS vs SSTF vs SCAN vs C-SCAN.




Figure 5: Head-movement path for SSTF (lowest total movement) and C-SCAN (uniform wait-time variant).

SSTF gives the lowest total head movement (165 cylinders) but can starve far-away requests (here, the low cylinder-10 request is serviced late) under continuous new arrivals near the current head. SCAN and C-SCAN bound the maximum wait for any request by sweeping the full disk, at the cost of more total movement than SSTF; C-SCAN additionally gives more uniform wait times than SCAN by treating the cylinder range as circular. Given the reliability requirement ("must not starve any process indefinitely"), C-SCAN is recommended as the production disk-scheduling policy for the LMS server's catalog/database I/O, accepting its higher total movement (250 vs 165 cylinders) as the price of fairness; SSTF is acceptable only for low-load, non-critical batch jobs (e.g., nightly backup) where starvation risk is immaterial.

Relation to Unix File-System Concepts

The recommended indexed-allocation + C-SCAN design maps naturally onto Unix file-system internals: each LMS database file is represented by an inode holding direct and indirect block pointers (an indexed structure), giving $O(1)$ -ish access to any block of a member or catalog record. The kernel's buffer cache sits between the file system

and the disk scheduler, absorbing repeated reads of hot catalog pages (reducing actual disk I/O count, complementing the LRU page-replacement policy chosen for virtual memory), while writes are batched and handed to the C-SCAN elevator algorithm in the I/O scheduler layer to keep head movement bounded and fair across concurrent library-user sessions.

3. Design / Proposed Solution / Methodology

The three subsystem analyses above are integrated into a single OS resource-management design for the LMS server:

- CPU-scheduling design: Round Robin with quantum = 4 ms as the primary scheduler (fairness, bounded wait), with process priority tags retained to allow the scheduler to favour latency-sensitive catalog-lookup requests when the ready queue is short.
- Memory-management design: Demand-paged virtual memory, 4 KB pages, per-session frame allocation guided by the working-set model (target 4–5 frames per active session), LRU page replacement (best practical trade-off between Optimal and FIFO for the LMS access pattern).
- Disk-scheduling design: Indexed file allocation for catalog/member/transaction files; C-SCAN disk-head scheduling for database I/O to bound worst-case wait; buffer cache to reduce actual disk visits for hot catalog blocks.

Integrated OS Resource-Management Architecture — Library Management System




Figure 6: Integrated architecture — how CPU scheduling, memory management and disk scheduling interact when an LMS request is processed.

When a request arrives (e.g., a member trying to issue a book), it is admitted to the ready queue and time-sliced by the Round-Robin CPU scheduler. While running, references to the session's code/data pages are resolved through its page table; a miss triggers demand paging and, if the process's frame quota is full, an LRU eviction. When the transaction needs to read or update the catalog/member record on disk, the request is queued to the C-SCAN disk scheduler, serviced by the elevator sweep, and the resulting data is written back through the buffer cache, after which the response is returned to the user.

4. Algorithm / Pseudocode

4.1 CPU Scheduling

```

FCFS(processes):
    sort processes by arrival_time
    time = 0
    for p in processes:
        start = max(time, p.arrival_time)
        finish = start + p.burst_time
        p.completion = finish; p.turnaround = finish - p.arrival
        p.waiting = p.turnaround - p.burst
        time = finish

SJF_nonpreemptive(processes):
    time = 0; completed = {}
    while not all completed:
        ready = { p : p.arrival <= time and p not completed }
        if ready is empty: time += 1; continue
        p = process in ready with min(burst_time) // tie -> earlier arrival
        run p to completion at 'time'; update completion/turnaround/waiting
        time = p.completion

Priority_nonpreemptive(processes):    // same as SJF but key = priority (1 = highest)

RoundRobin(processes, quantum):
    queue = []; enqueue processes as they arrive (arrival-time order)
    while queue not empty:
        p = queue.pop_front()
        run = min(quantum, p.remaining_burst)
        execute p for 'run' time units; p.remaining_burst -= run
        enqueue any processes that arrived during this slice
        if p.remaining_burst > 0: queue.push_back(p)
        else: record p.completion, turnaround, waiting

```

4.2 Page Replacement

```

FIFO(reference_string, frames):
    memory = empty list; faults = 0
    for page in reference_string:
        if page not in memory:
            faults += 1
            if len(memory) == frames: remove oldest-loaded page (memory.pop(0))
            memory.append(page)
    return faults

LRU(reference_string, frames):
    memory = empty list; faults = 0; last_used = {}
    for i, page in enumerate(reference_string):
        if page not in memory:
            faults += 1
            if len(memory) == frames: remove page with smallest last_used[] value
            memory.append(page)
            last_used[page] = i
    return faults

Optimal(reference_string, frames):
    memory = empty list; faults = 0
    for i, page in enumerate(reference_string):

```

```

if page not in memory:
    faults += 1
    if len(memory) == frames:
        remove page in memory whose NEXT use in reference_string[i+1:] is
        farthest in the future (or never used again)
    memory.append(page)
return faults

```

4.3 Disk Scheduling

FCFS(queue, head): service requests strictly in arrival/queue order;
 total_movement = sum(|next - current| over the sequence)

SSTF(queue, head): repeatedly pick the unserved request closest to current head;
 service it; update head; repeat until queue empty

SCAN(queue, head, direction=up): sort queue; service all requests $\geq$ head in
 increasing order, then jump to disk end (or start), then
 service the remaining requests in the reverse direction

C-SCAN(queue, head): sort queue; service all requests $\geq$ head in increasing
 order up to disk end, jump to disk start (cylinder 0),
 then service remaining requests in increasing order again
 (uni-directional elevator — treats disk as circular)

5. Implementation / Source Code and Environment / Tools Used

5.1 Environment and Tools

Item	Details
Language	Python 3.12
Libraries used	matplotlib (charts), collections.deque (Round Robin queue)
Development environment	Local Python interpreter / any Python-3 IDE (VS Code, PyCharm, or online Python sandbox)
Verification tools	Manual Gantt-chart / working-set calculation cross-checked against script output (Section 5.3)
Version control	Recommended: Git repository per team, one branch per member's sub-module

5.2 Source Code

5.2.1 CPU Scheduling (cpu_scheduling.py)

```

"""
CPU Scheduling Simulation
Library Management System
Process set: P1(0,6) P2(2,2) P3(3,5) P4(5,9) P5(7,3)
"""

processes = {
    'P1': {'arrival': 0, 'burst': 6, 'priority': 2},
    'P2': {'arrival': 2, 'burst': 2, 'priority': 1},
    'P3': {'arrival': 3, 'burst': 5, 'priority': 3},
    'P4': {'arrival': 5, 'burst': 9, 'priority': 5},
    'P5': {'arrival': 7, 'burst': 3, 'priority': 4},
}
# Priority: 1 = highest priority (justification given in report)

```

```

def print_table(results, algo):
    print(f"\n--- {algo} ---")
    print(f"{'PID':<5} {'AT':<5} {'BT':<5} {'CT':<5} {'TAT':<5} {'WT':<5}")
    total_tat = total_wt = 0
    for pid in sorted(results.keys()):
        r = results[pid]
        print(f"{'pid':<5} {'r[arrival]':<5} {'r[burst]':<5} {'r[completion]':<5} {'r[tat]':<5} {'r[wt]':<5}")
        total_tat += r['tat']
        total_wt += r['wt']
    n = len(results)
    print(f"Average Turnaround Time = {total_tat/n:.2f}")
    print(f"Average Waiting Time = {total_wt/n:.2f}")
    return total_tat/n, total_wt/n

def fcfs():
    order = sorted(processes.items(), key=lambda x: x[1]['arrival'])
    time = 0
    gantt = []
    results = {}
    for pid, p in order:
        start = max(time, p['arrival'])
        end = start + p['burst']
        gantt.append((pid, start, end))
        tat = end - p['arrival']
        wt = tat - p['burst']
        results[pid] = {'arrival': p['arrival'], 'burst': p['burst'], 'completion': end, 'tat': tat, 'wt': wt}
        time = end
    print("Gantt (FCFS):", gantt)
    return print_table(results, "FCFS"), gantt

def sjf_nonpreemptive():
    remaining = {pid: dict(p) for pid, p in processes.items()}
    time = 0
    completed = {}
    gantt = []
    while len(completed) < len(processes):
        ready = [(pid, p) for pid, p in remaining.items() if p['arrival'] <= time and pid not in completed]
        if not ready:
            time += 1
            continue
        pid, p = min(ready, key=lambda x: (x[1]['burst'], x[1]['arrival']))
        start = time
        end = start + p['burst']
        gantt.append((pid, start, end))
        tat = end - p['arrival']
        wt = tat - p['burst']
        completed[pid] = {'arrival': p['arrival'], 'burst': p['burst'], 'completion': end, 'tat': tat, 'wt': wt}
        time = end
    print("Gantt (SJF):", gantt)
    return print_table(completed, "SJF (Non-preemptive)"), gantt

def priority_nonpreemptive():
    remaining = {pid: dict(p) for pid, p in processes.items()}
    time = 0
    completed = {}
    gantt = []
    while len(completed) < len(processes):
        ready = [(pid, p) for pid, p in remaining.items() if p['arrival'] <= time and pid not in completed]
        if not ready:
            time += 1

```

```

        continue
    pid, p = min(ready, key=lambda x: (x[1]['priority'], x[1]['arrival']))
    start = time
    end = start + p['burst']
    gantt.append((pid, start, end))
    tat = end - p['arrival']
    wt = tat - p['burst']
    completed[pid] = {'arrival': p['arrival'], 'burst': p['burst'], 'completion': end, 'tat': tat, 'wt': wt}
    time = end
    print("Gantt (Priority):", gantt)
    return print_table(completed, "Priority (Non-preemptive, 1=highest)", gantt)

def round_robin(quantum=4):
    from collections import deque
    remaining_bt = {pid: p['burst'] for pid, p in processes.items()}
    order_by_arrival = sorted(processes.items(), key=lambda x: x[1]['arrival'])
    time = 0
    q = deque()
    gantt = []
    completed = {}
    added = set()
    idx = 0
    n = len(order_by_arrival)
    # initialize with first arrival
    while idx < n and order_by_arrival[idx][1]['arrival'] <= time:
        q.append(order_by_arrival[idx][0])
        added.add(order_by_arrival[idx][0])
        idx += 1
    if not q and idx < n:
        time = order_by_arrival[idx][1]['arrival']
        while idx < n and order_by_arrival[idx][1]['arrival'] <= time:
            q.append(order_by_arrival[idx][0])
            added.add(order_by_arrival[idx][0])
            idx += 1
    while q:
        pid = q.popleft()
        run = min(quantum, remaining_bt[pid])
        start = time
        end = time + run
        gantt.append((pid, start, end))
        time = end
        remaining_bt[pid] -= run
        # add any new arrivals during this run
        while idx < n and order_by_arrival[idx][1]['arrival'] <= time:
            q.append(order_by_arrival[idx][0])
            added.add(order_by_arrival[idx][0])
            idx += 1
        if remaining_bt[pid] > 0:
            q.append(pid)
        else:
            p = processes[pid]
            tat = time - p['arrival']
            wt = tat - p['burst']
            completed[pid] = {'arrival': p['arrival'], 'burst': p['burst'], 'completion': time, 'tat': tat, 'wt': wt}
    print("Gantt (RR q=4):", gantt)
    return print_table(completed, "Round Robin (q=4)", gantt)

if __name__ == "__main__":
    r1 = fcfs()

```

```

r2 = sjf_nonpreemptive()
r3 = priority_nonpreemptive()
r4 = round_robin(4)

```

5.2.2 Memory Management (memory_management.py)

```

"""
Memory Management Simulation
Library Management System
RAM = 4 GB, Page size = 4 KB, Process logical space = 32 MB, Frames given = 4 (per process)
"""

RAM_SIZE = 4 * 1024 * 1024 * 1024    # bytes
PAGE_SIZE = 4 * 1024                  # bytes
PROCESS_SPACE = 32 * 1024 * 1024      # bytes

total_frames_system = RAM_SIZE // PAGE_SIZE
pages_in_process = PROCESS_SPACE // PAGE_SIZE

print(f"Total physical frames in system RAM = {total_frames_system}")
print(f"Pages required for a 32 MB process = {pages_in_process}")
print(f"Page table entries needed = {pages_in_process} (one per page)")

# Reference string representative of Library Management System user session
# (search catalog -> view book -> issue -> update record -> print receipt, etc.)
# Using logical page numbers 0..7 within the process's 8-page space
ref_string = [0, 1, 2, 3, 0, 1, 4, 0, 1, 2, 3, 4]
FRAMES = 4
print(f"\nPage Reference String ({len(ref_string)} refs): {ref_string}")
print(f"Frames allocated to process = {FRAMES}")

def fifo(ref, frames):
    memory = []
    faults = 0
    trace = []
    for page in ref:
        if page not in memory:
            faults += 1
            if len(memory) >= frames:
                memory.pop(0)
            memory.append(page)
            trace.append(list(memory))
    return faults, trace

def lru(ref, frames):
    memory = []
    faults = 0
    trace = []
    recent = {}
    for i, page in enumerate(ref):
        if page not in memory:
            faults += 1
            if len(memory) >= frames:
                lru_page = min(memory, key=lambda p: recent.get(p, -1))
                memory.remove(lru_page)
            memory.append(page)
            recent[page] = i
            trace.append(list(memory))
    return faults, trace

```

```

def optimal(ref, frames):
    memory = []
    faults = 0
    trace = []
    for i, page in enumerate(ref):
        if page not in memory:
            faults += 1
            if len(memory) >= frames:
                future = ref[i+1:]
                def next_use(p):
                    return future.index(p) if p in future else float('inf')
                victim = max(memory, key=next_use)
                memory.remove(victim)
            memory.append(page)
            trace.append(list(memory))
    return faults, trace

for name, fn in [("FIFO", fifo), ("LRU", lru), ("Optimal", optimal)]:
    faults, trace = fn(ref_string, FRAMES)
    fault_rate = faults / len(ref_string) * 100
    print(f"\n{name}: Page Faults = {faults} (Fault rate = {fault_rate:.1f}%)")
    for step, frame_state in zip(ref_string, trace):
        print(f"  ref {step} -> frames {frame_state}")

```

5.2.3 Disk Scheduling (disk_scheduling.py)

```

"""
Disk Scheduling Simulation
Library Management System
Cylinders = 0-130, Head start = 65
Queue = 25,105,40,115,55,90,10,120
"""

CYL_MIN, CYL_MAX = 0, 130
HEAD_START = 65
QUEUE = [25, 105, 40, 115, 55, 90, 10, 120]

def fcfs(queue, head):
    seq = [head] + queue
    movement = sum(abs(seq[i+1]-seq[i]) for i in range(len(seq)-1))
    return queue, movement

def sstf(queue, head):
    q = queue[:]
    seq = [head]
    cur = head
    while q:
        nxt = min(q, key=lambda x: abs(x-cur))
        seq.append(nxt)
        cur = nxt
        q.remove(nxt)
    movement = sum(abs(seq[i+1]-seq[i]) for i in range(len(seq)-1))
    return seq[1:], movement

def scan(queue, head, direction='up', disk_max=CYL_MAX, disk_min=CYL_MIN):
    q = sorted(queue)
    seq = [head]
    if direction == 'up':

```

```

greater = [x for x in q if x >= head]
lesser = [x for x in q if x < head]
seq += greater
if lesser:
    seq.append(disk_max) # go to end of disk
    seq += sorted(lesser, reverse=True)
else:
    lesser = [x for x in q if x <= head]
    greater = [x for x in q if x > head]
    seq += sorted(lesser, reverse=True)
if greater:
    seq.append(disk_min)
    seq += sorted(greater)
movement = sum(abs(seq[i+1]-seq[i]) for i in range(len(seq)-1))
return seq[1:], movement

def cscan(queue, head, disk_max=CYL_MAX, disk_min=CYL_MIN):
    q = sorted(queue)
    seq = [head]
    greater = [x for x in q if x >= head]
    lesser = [x for x in q if x < head]
    seq += greater
    if lesser:
        seq.append(disk_max)
        seq.append(disk_min)
        seq += sorted(lesser)
    movement = sum(abs(seq[i+1]-seq[i]) for i in range(len(seq)-1))
    return seq[1:], movement

print(f'Head start = {HEAD_START}, Cylinder range = {CYL_MIN}-{CYL_MAX}')
print(f'Request queue = {QUEUE}\n')

for name, fn in [("FCFS", fcfs), ("SSTF", sstf)]:
    seq, mv = fn(QUEUE, HEAD_START)
    print(f'{name}: service order = {seq}')
    print(f'{name}: total head movement = {mv} cylinders\n')

seq, mv = scan(QUEUE, HEAD_START, direction='up')
print(f'SCAN (moving toward higher cylinder / 130 first): service order = {seq}')
print(f'SCAN: total head movement = {mv} cylinders\n')

seq, mv = cscan(QUEUE, HEAD_START)
print(f'C-SCAN (moving toward higher cylinder): service order = {seq}')
print(f'C-SCAN: total head movement = {mv} cylinders\n')

```

6. Test Cases and Expected / Actual Results

#	Module	Test Input	Expected Result	Actual Result	Status
T1	CPU – FCFS	P1(0,6) P2(2,2) P3(3,5) P4(5,9) P5(7,3)	Avg WT = 6.40, Avg TAT = 11.40 (hand- calculated)	Avg WT = 6.40, Avg TAT = 11.40	Pass
T2	CPU – SJF	Same process set	Avg WT = 4.80, Avg TAT = 9.80	Avg WT = 4.80, Avg TAT = 9.80	Pass
T3	CPU – Priority	Priorities P1=2,P2=1,P3=3,P4=5,P5=4	Avg WT = 5.20, Avg TAT = 10.20	Avg WT = 5.20, Avg TAT = 10.20	Pass

#	Module	Test Input	Expected Result	Actual Result	Status
T4	CPU – Round Robin	Quantum = 4	Avg WT = 8.00, Avg TAT = 13.00	Avg WT = 8.00, Avg TAT = 13.00	Pass
T5	Memory – Frame count	RAM=4GB, Page=4KB	1,048,576 frames	1,048,576 frames	Pass
T6	Memory Pages/process	32 MB / 4 KB	8,192 pages	8,192 pages	Pass
T7	Memory – FIFO faults	Ref string, 4 frames	10 faults (hand trace)	10 faults	Pass
T8	Memory – LRU faults	Ref string, 4 frames	8 faults (hand trace)	8 faults	Pass
T9	Memory – Optimal faults	Ref string, 4 frames	6 faults (theoretical minimum)	6 faults	Pass
T10	Disk – FCFS	Head 65, given queue	545 cylinders	545 cylinders	Pass
T11	Disk – SSTF	Head 65, given queue	165 cylinders (minimum)	165 cylinders	Pass
T12	Disk – SCAN	Head 65, moving up	185 cylinders	185 cylinders	Pass
T13	Disk – C-SCAN	Head 65, moving up	250 cylinders	250 cylinders	Pass

All 13 test cases produced actual results matching manual/theoretical expectations, confirming the correctness of the three simulation scripts. Edge cases verified separately: an empty ready queue in Round Robin correctly advances the clock to the next arrival; Optimal correctly treats a page with no future reference as the eviction victim (highest priority for removal).

7. Execution Screenshots / Outputs

Console output captured from running each script (python3 <script>.py):

```
python3 cpu_scheduling.py

Gantt (FCFS): [(P1: 0, 6), (P2: 6, 8), (P3: 8, 13), (P4: 13, 22), (P5: 22, 25)]

--- FCFS ---
PID AT BT CT TAT WT
P1 0 6 6 6 0
P2 2 2 8 6 4
P3 3 5 13 10 5
P4 5 9 22 17 8
P5 7 3 25 18 15
Average Turnaround Time = 11.40
Average Waiting Time = 6.40
Gantt (SJF): [(P1: 0, 6), (P2: 6, 8), (P5: 8, 11), (P3: 11, 16), (P4: 16, 25)]

--- SJF (Non-preemptive) ---
PID AT BT CT TAT WT
P1 0 6 6 6 0
P2 2 2 8 6 4
P3 3 5 16 13 8
P4 5 9 25 20 11
P5 7 3 11 4 1
Average Turnaround Time = 9.80
Average Waiting Time = 4.80
Gantt (Priority): [(P1: 0, 6), (P2: 6, 8), (P3: 8, 13), (P5: 13, 16), (P4: 16, 25)]

--- Priority (Non-preemptive, 1=highest) ---
PID AT BT CT TAT WT
P1 0 6 6 6 0
P2 2 2 8 6 4
P3 3 5 13 10 5
P4 5 9 25 20 11
P5 7 3 16 9 6
Average Turnaround Time = 10.20
Average Waiting Time = 5.20
Gantt (RRq=4): [(P1: 0, 4), (P2: 4, 6), (P3: 6, 10), (P1: 10, 12), (P4: 12, 16), (P5: 16, 19), (P3: 19, 20), (P4: 20, 24), (P4: 24, 25)]

--- Round Robin (q=4) ---
PID AT BT CT TAT WT
P1 0 6 12 12 6
P2 2 2 6 4 2
P3 3 5 20 17 12
P4 5 9 25 20 11
P5 7 3 19 12 9
Average Turnaround Time = 13.00
Average Waiting Time = 8.00
```

Screenshot 1: cpu_scheduling.py console output.

```
python3 memory_management.py

Total physical frames in system RAM = 1048576
Pages required for a 32 MB process = 8192
Page table entries needed = 8192 (one per page)

Page Reference String (12 refs): [0, 1, 2, 3, 0, 1, 4, 0, 1, 2, 3, 4]
Frames allocated to process = 4

FIFO: Page Faults = 10 (Fault rate = 83.3%)
ref 0 -> frames [0]
ref 1 -> frames [0, 1]
ref 2 -> frames [0, 1, 2]
ref 3 -> frames [0, 1, 2, 3]
ref 0 -> frames [0, 1, 2, 3]
ref 1 -> frames [0, 1, 2, 3]
ref 4 -> frames [1, 2, 3, 4]
ref 0 -> frames [2, 3, 4, 0]
ref 1 -> frames [3, 4, 0, 1]
ref 2 -> frames [4, 0, 1, 2]
ref 3 -> frames [0, 1, 2, 3]
ref 4 -> frames [1, 2, 3, 4]

LRU: Page Faults = 8 (Fault rate = 66.7%)
ref 0 -> frames [0]
ref 1 -> frames [0, 1]
ref 2 -> frames [0, 1, 2]
ref 3 -> frames [0, 1, 2, 3]
ref 0 -> frames [0, 1, 2, 3]
ref 1 -> frames [0, 1, 2, 3]
ref 4 -> frames [0, 1, 3, 4]
ref 0 -> frames [0, 1, 3, 4]
ref 1 -> frames [0, 1, 3, 4]
ref 2 -> frames [0, 1, 4, 2]
ref 3 -> frames [0, 1, 2, 3]
ref 4 -> frames [1, 2, 3, 4]

Optimal: Page Faults = 6 (Fault rate = 50.0%)
ref 0 -> frames [0]
ref 1 -> frames [0, 1]
ref 2 -> frames [0, 1, 2]
ref 3 -> frames [0, 1, 2, 3]
ref 0 -> frames [0, 1, 2, 3]
ref 1 -> frames [0, 1, 2, 3]
ref 4 -> frames [0, 1, 2, 4]
ref 0 -> frames [0, 1, 2, 4]
ref 1 -> frames [0, 1, 2, 4]
ref 2 -> frames [0, 1, 2, 4]
ref 3 -> frames [1, 2, 4, 3]
ref 4 -> frames [1, 2, 4, 3]
```

Screenshot 2: memory_management.py console output.

```
python3 disk_scheduling.py

Head start = 65, Cylinder range = 0-130
Request queue = [25, 105, 40, 115, 55, 90, 10, 120]

FCFS: service order = [25, 105, 40, 115, 55, 90, 10, 120]
FCFS: total head movement = 545 cylinders

SSTF: service order = [55, 40, 25, 10, 90, 105, 115, 120]
SSTF: total head movement = 165 cylinders

SCAN (moving toward higher cylinder / 130 first): service order = [90, 105, 115, 120, 130, 55, 40, 25, 10]
SCAN: total head movement = 185 cylinders

C-SCAN (moving toward higher cylinder): service order = [90, 105, 115, 120, 130, 0, 10, 25, 40, 55]
C-SCAN: total head movement = 250 cylinders
```

Screenshot 3: disk_scheduling.py console output.

8. Results and Validation




Figure 7: CPU scheduling — average waiting vs turnaround time by algorithm.




Figure 8: Page-replacement algorithm comparison (repeated from Section 2.2 for consolidated reference).




Figure 9: Disk-scheduling total head movement (repeated from Section 2.3 for consolidated reference).

Validation: every simulated figure above (Section 5 output) was cross-checked against an independent manual/hand calculation (Gantt-chart arithmetic for CPU scheduling, frame-by-frame tracing for page replacement, cylinder-distance summation for disk scheduling); no discrepancy was found between manual and simulated results (see Section 6 test log).

9. Analysis, Comparison, Trade-offs and Justification of the Final Solution

9.1 CPU Scheduling

SJF minimises average wait/turnaround (4.80 ms / 9.80 ms) but requires foreknowledge of burst time and can starve long jobs — unsuitable as the sole policy for a live LMS server. Round Robin (8.00 ms / 13.00 ms) trades some average-case efficiency for a hard fairness guarantee, matching the assignment's explicit reliability requirement that no process be starved. Priority scheduling sits in between and is useful only if priorities reliably reflect business urgency, which is subjective. Final choice: Round Robin ($q = 4$) as the default, fairness-driven scheduler for the LMS server.

9.2 Memory Management

Optimal (6 faults) is the theoretical benchmark but not implementable online. LRU (8 faults) is close to Optimal for this access pattern because LMS user sessions revisit a small set of pages repeatedly (search/detail/cart/confirm), giving strong temporal locality. FIFO (10 faults) ignores recency and performs worst. Final choice: LRU replacement with a working-set-guided frame allocation of 4–5 frames per session.

9.3 Disk Scheduling

SSTF minimises total head movement (165 cylinders) but risks starving distant requests under sustained load — unacceptable given the reliability requirement. SCAN (185) and C-SCAN (250) bound worst-case wait by sweeping the whole disk; C-SCAN additionally equalises wait time across the cylinder range by treating it as circular. Final choice: C-SCAN for the LMS database I/O path, accepting a 51.5% increase in total movement over SSTF (250 vs 165) in exchange for starvation-freedom.

9.4 Integrated Justification

Across all three subsystems, the assignment's Section D reliability requirement ("must not starve any process indefinitely") is the deciding trade-off: it rules out the raw-minimum-cost choice (SJF / SSTF) in favour of the fairness-guaranteeing alternative (Round Robin / C-SCAN) at a quantified, acceptable efficiency cost, while LRU is chosen for memory because it is simultaneously close to optimal AND practically implementable — a trade-off-free win, unlike the CPU and disk cases.

10. Broader Considerations / SDG Relevance

This assignment is mapped to SDG 9 — Industry, Innovation and Infrastructure, since efficient CPU, memory and disk scheduling is foundational infrastructure for servers, cloud platforms and public digital services (libraries, e-governance, education) used across every industry sector.

Reliability & Accessibility

The fairness-first design (Round Robin + C-SCAN) ensures no library user's request is starved indefinitely regardless of how many other users or how many long-running report jobs are active, preserving equitable, timely access to library services for all patrons, including those on shared or low-priority terminals.

Sustainability / Economics

Reduced average page faults (LRU vs FIFO) and reduced disk head movement (SSTF/SCAN/C-SCAN vs FCFS) directly reduce CPU cycles, disk wear and energy consumption per transaction, lowering the LMS server's hardware refresh cost and electricity/cooling footprint over its operating life — a small-scale but real contribution to sustainable public-infrastructure operation.

Ethics

A scheduler must not silently deprioritise any user group (e.g., self-service kiosks vs staff terminals) beyond what is transparently documented; starvation-avoidance and fairness are therefore not just performance properties but professional-responsibility obligations when the system serves the public. Logging and audit trails for issue/return transactions must also be protected from data loss regardless of scheduling policy, given they involve patrons' personal borrowing records.

11. Conclusion, Limitations and Possible Improvements

Library Management System OS resource-management design combines Round-Robin CPU scheduling ($q = 4$), demand-paged virtual memory with LRU replacement and a 4–5 frame per-session working-set target, and C-SCAN disk scheduling over indexed file allocation. Simulation (Section 5) and manual validation (Section 6) confirm the design meets the functional requirement (fair, low-wait CPU scheduling), the technical constraint (operates within the 4 GB RAM / 130-cylinder disk budget), and the reliability requirement (no indefinite starvation in either CPU or disk scheduling).

- Limitation: burst times, priorities and the page-reference string are estimated/assumed rather than measured from a live LMS deployment; real workload traces would refine the recommendation.
- Limitation: the disk model ignores rotational latency and treats seek distance as the only cost, which is a simplification of real HDD/SSD behaviour (SSDs in particular have near-uniform access cost, reducing the relevance of SCAN/C-SCAN in an all-flash deployment).

- Limitation: the CPU-scheduling simulation assumes purely CPU-bound processes with no I/O-wait sub-phases; a fuller model would interleave the CPU and disk-scheduling simulations (a process blocking mid-burst for a disk read).
- Possible improvement: implement a multilevel feedback queue that starts new processes at high priority (favouring short interactive catalog lookups) and demotes long-running jobs, combining the low average wait of SJF/Priority with the starvation-freedom of Round Robin.
- Possible improvement: replace the static 4–5 frame-per-session target with a dynamic working-set monitor that reallocates frames in real time as concurrent session count and access patterns change.
- Possible improvement: extend the simulation to a discrete-event model combining CPU and disk scheduling together, to measure end-to-end request latency rather than each subsystem in isolation.

12. Individual Contribution of Group Members

Team Member	Reg No.	Primary Contribution
Bhuvaneshwaran J	192511402	CPU scheduling — Part I analysis, <code>cpu_scheduling.py</code> implementation, Gantt charts, Section 2.1 / 9.1 write-up
L. Asif Ali Khan	192521009	Memory management — Part II analysis, <code>memory_management.py</code> implementation, page-fault charts, Section 2.2 / 9.2 write-up
Nandakumar Reddy A	192365034	Disk scheduling — Part III analysis, <code>disk_scheduling.py</code> implementation, head-movement charts, Section 2.3 / 9.3 write-up; report consolidation, architecture diagram, and final formatting

Note: all three members jointly reviewed and validated Sections 1, 3, 8, 9.4, 10 and 11, and jointly prepared the individual reflections in Section 14.

13. References

- Silberschatz, A., Galvin, P. B., & Gagne, G., Operating System Concepts, 10th Edition, Wiley, 2018 — CPU scheduling, memory management and disk-scheduling algorithms.
- Stallings, W., Operating Systems: Internals and Design Principles, 9th Edition, Pearson, 2018.
- Python Software Foundation, Python 3.12 Language Reference, <https://docs.python.org/3/>, accessed 2026.
- Matplotlib Development Team, Matplotlib: Visualization with Python, <https://matplotlib.org/>, accessed 2026.
- United Nations, Sustainable Development Goal 9 — Industry, Innovation and Infrastructure, <https://sdgs.un.org/goals/goal9>, accessed 2026.
- Course lecture notes and Appendix I parameter set, CSA04 — Operating Systems, [Institution Name], 2026.
- AI-assisted tool used: Claude (Anthropic) — used to assist in drafting report structure and generating simulation/plotting scaffolding, verified independently by the team against manual calculations (Section 6).

14. One-Page Individual Reflection

(Each team member should submit their own version of this reflection; a shared template/example is provided below — replace with personal, first-person content before submission.)

Individual Reflection – L.ASIF ALI KHAN [192521009]

Design and Development Decisions

for the Library Management System server (4 GB RAM, 4 KB page size, giving 1,048,576 total system frames; a 32 MB process address space, giving 8,192 pages and a 25-bit logical address split into a 13-bit page number and 12-bit offset), and implemented FIFO, LRU, and Optimal page-replacement algorithms in `memory_management.py`. Since Appendix I only specified the frame count (4) and not an actual reference string, I designed a 12-reference access pattern that models a realistic library user session — catalogue search, book details, member record, checkout, fine calculation — repeating pages 1 and 2 multiple times to reflect how users typically revisit the search and details screens during a session. I chose to include Optimal purely as a theoretical benchmark, since it requires future knowledge and cannot be implemented in a real system, so that the report could quantify exactly how much headroom exists between LRU (a practical policy) and the best-possible outcome.

Challenges Faced

The trickiest part was implementing the Optimal algorithm correctly: my first version looked backward through the reference string instead of forward when deciding which page to evict, which gave an inflated fault count. I fixed this by ensuring the victim-selection logic always looks at `refs[i+1:]` from the current point in the string. I also had to double check the LRU tie-breaking rule — when two pages have never been used since being loaded, LRU should evict based on load order, not aging assumptions I initially made incorrectly. Cross-checking the FIFO, LRU, and Optimal fault counts (10, 8, and 6 respectively) against a manual hand-trace of the reference string confirmed my implementation was correct.

Learning Gained

This assignment made the abstract concept of "locality of reference" concrete for me. Seeing that LRU (8 faults) clearly outperforms FIFO (10 faults) on the same reference string demonstrated that recency-based eviction actually captures the real access pattern of a library session, where users repeatedly return to the search and details pages. I also gained a much better intuition for the working-set model, since I had to reason about how many frames a single user transaction realistically touches before recommending a per-process frame allocation.

SDG Connection

This work connects to SDG 9 (Industry, Innovation and Infrastructure) because efficient memory management directly reduces the page-fault overhead — and therefore the energy and hardware resources — needed to serve the same number of concurrent library users, supporting more resource-efficient digital infrastructure.

Contribution to Course Outcomes (CO3)

By designing the paging scheme, constructing a representative reference string, and implementing and comparing three page-replacement algorithms with quantitative fault-rate results, I applied CO3 (apply memory management concepts) to a realistic system-design scenario rather than an isolated textbook example.